

# Township Clerk AI Assistant – Quick Guide

AnythingLLM Desktop + eGPU (RTX 4000) + External Drive Models (Samsung T5). Print-friendly cheat sheets for daily use.

## Quick “Do I need the external drive plugged in?”

- If your model files (.gguf) live on the Samsung T5, the T5 must be plugged in to use those models.
- If you imported/copied models into the laptop storage, the T5 is optional for model use (but still useful for documents/backups).

## 1) Which Model to Use (Cheat Sheet)

Model	Best for	Avoid when / Notes
<b>Llama 3.1 8B (Instruct)</b> “Everyday Assistant”	<ul style="list-style-type: none"><li>• Drafting emails/letters, polite notices, public-facing wording</li><li>• Summaries of typical documents</li><li>• General Q&amp;A (policies, procedures, “explain this to me”)</li></ul>	<ul style="list-style-type: none"><li>• If it feels too verbose, ask: “Answer in 5 bullets, no fluff.”</li><li>• For very long documents, consider Phi-3 Mini 128k.</li></ul>
<b>Mistral 7B (Instruct)</b> “Logic & Formal Tone”	<ul style="list-style-type: none"><li>• More formal, structured answers</li><li>• Step-by-step reasoning (“What should I do next?”)</li><li>• Comparisons (“Option A vs B”) and checklists</li></ul>	<ul style="list-style-type: none"><li>• If you need maximum context length (very large text), Phi-3 Mini may handle big inputs better.</li></ul>
<b>Phi-3 Mini 128k (Instruct)</b> “Long Document Reader”	<ul style="list-style-type: none"><li>• Long meeting minutes, large PDFs, big ordinance packs</li><li>• “Find the section that says...” and “What changed between versions?”</li><li>• Pulling details across many pages</li></ul>	<ul style="list-style-type: none"><li>• Sometimes less “chatty”/creative; best used for document-heavy tasks.</li><li>• If you only need a quick email draft, Llama is usually faster/easier.</li></ul>

#### **Fast pick (if you're unsure):**

- Writing/drafting/general help → **Llama 3.1 8B**
- Formal, structured checklists & decisions → **Mistral 7B**
- Very long documents / “find it in the PDF” → **Phi-3 Mini 128k**

## **2) Daily Setup Procedure (Start of Day)**

### **A) Hardware (2 minutes)**

- Turn on the eGPU enclosure (Razer Core) and confirm its power/fans are on.
- Connect the Thunderbolt/USB-C cable from the eGPU to the laptop (lightning-bolt USB-C port).
- Plug in the Samsung T5 (if your models are stored on it).

### **B) Windows quick check (optional but recommended)**

- Open **Task Manager** → **Performance** tab → verify an NVIDIA GPU graph exists (this confirms the eGPU is recognized).

### **C) AnythingLLM (1 minute)**

- Open **AnythingLLM Desktop**.
- Select your workspace (example: **Township Assistant**).
- Ask a quick test question: “Summarize today’s priorities in 5 bullets.”

## **3) How to Change Models (Cheat Sheet)**

### **Method 1 (common): Change the model in LLM Preference**

- In AnythingLLM, open **Settings**.
- Go to **LLM Preference** (or the LLM/Provider screen).
- Select the model you want (Llama / Mistral / Phi).
- Return to your workspace chat and continue.

### **Method 2: Import models (only if you don’t see them yet)**

- Settings → **LLM Preference**.
- Click **Import GGUF file**.
- Browse to the Samsung T5 and select the .gguf file.
- Repeat for each model you want available.

#### **Common gotcha:**

- If AnythingLLM can’t find a model you previously used, the Samsung T5 may be unplugged or the drive letter changed.
- If a model loads but responses are slow, confirm the eGPU is connected and recognized in Windows.

#### 4) Best “Ask Prompts” for Clerk Work

Task	Copy/paste prompt
Draft a professional email	Draft a professional email reply. Tone: polite, concise, public-sector. Include next steps as bullets. Here are the details: ...
Summarize minutes	Summarize the key decisions, motions, votes, and follow-up actions. Output as: (1) Decisions (2) Action items (3) Deadlines (4) Open questions.
Create a checklist	Turn this policy into a step-by-step checklist a clerk can follow. Keep it to 10 bullets max.
Explain in plain language	Explain this section in plain English for a resident. Keep it friendly and under 150 words.
Find a detail in a long PDF	Search the attached document(s) and quote the exact sentence about: [topic]. Then summarize what it means in 3 bullets.

#### 5) Shutdown Procedure (End of Day)

##### A) Close software cleanly

- In AnythingLLM, finish the current response (don’t unplug while it’s “thinking”).
- Close AnythingLLM (File → Exit or click the X).

##### B) Eject external drives

- In Windows, click the USB “Safely Remove Hardware” icon (system tray) and eject the Samsung T5.

##### C) Disconnect eGPU

- Unplug the Thunderbolt/USB-C cable from the laptop.
- Turn off the eGPU enclosure (if you won’t use it again soon).

##### Two safety rules:

- Don’t unplug the eGPU or Samsung T5 while AnythingLLM is actively generating a response.
- If anything acts weird, do the classic IT fix: close AnythingLLM → disconnect eGPU → reconnect eGPU → reopen AnythingLLM.

## 6) Quick Troubleshooting

- **Models disappeared:** Plug Samsung T5 back in and re-open AnythingLLM. If needed, re-import via “Import GGUF file.”
- **It’s slow:** Check eGPU connection, confirm NVIDIA driver installed, and look at Task Manager → Performance → NVIDIA GPU activity during a response.
- **No microphone / voice:** Windows Settings → Privacy & security → Microphone → allow access for AnythingLLM.

Version 1.0 • Built for Dave’s “Township Clerk Assistant” setup (AnythingLLM + RTX 4000 eGPU + Samsung T5 models). Print on A4 or Letter.